

VideoSQL: Ultra Long Video Understanding by Agentic SQL Queries with Trainable Skills

Anonymous CVPR submission

Paper ID *****

Abstract

1. Introduction

Ultra long video understanding (ULVU)

2. Related Works

2.1. Ultra Long Video Understanding

2.2. SQL-based RAG

Relational databases have long been the default for storing structured data. Recent work starts applying them to agentic RAG and agent memory. SAG [11] argues that naive RAG yields shallow retrieval, while graph-based RAG is slow, costly, and hard to update. They therefore use SQL as the backbone of RAG. Documents are stored as an event-entity incidence table. At query time, an LLM extracts entities from the query, and a JOIN retrieves events sharing these entities as seeds. A reverse JOIN then collects the entities attached to these seed events, and a further JOIN brings in previously unseen events incident to them. The expanded event set is passed back to the LLM, which selects the evidence and maps it to the original chunks. There are also other works that frame agent memories as databases like [7].

2.3. Skill Training

Training model weights, especially with reinforcement learning (RL), is expensive, slow, and not portable (you would need to host and run this specific model, which might perform poorly on general tasks). Many recent works began to train the agent's Skill document instead of model weights, which can be seen as a form of prompt optimization.

Training-free GRPO [1] uses an approach similar to the original gradient-based GRPO [9]. An initial experience and a policy model are used to roll out a batch of traces, and another LLM computes the semantic advantage of each

rollout, analogous to the numeric advantage in the original GRPO. These semantic advantages are used to update the experience. **Trace2Skill** [6] uses parallel rollouts in batches to collect many traces and generate patches for successful and failed traces. These patches are then merged into the final Skill in sequence. Compared to training-free GRPO, this can be seen as a form of offline reinforcement learning. **SkillOpt** [13] treats Skill optimization similarly to gradient descent, where the editing budget is analogous to the learning rate, the gradient direction is analogous to the editing direction, etc. It also learns a meta-skill to track what kind of edit is helpful using a held-out validation set. **SkillCAT** [3] uses DPO-style training. It samples two paired rollouts: a successful one (r^+) and a failed one (r^-). It then analyzes the two traces, finds the point where they diverge, and generates a patch contrastively. The patch is then verified to see whether it helps the agent, similar to the gate in SkillOpt, and is merged if it turns a failure into a success or maintains a success. However, in many cases, the two traces are completely independent (with no shared prefix beyond the prompt), or the divergence point does not affect the result.

3. VideoSQL

3.1. Indexing

For a fair comparison with our baseline MAGIC-Video [4], we re-used their index. We convert their index from a graph data structure into a relational database by constructing a table for each node type, and use an outer key as an edge. We also formulated the topic chain and event chain into tables. This yields major tables of episode captions, semantic triplets, entities, and topic and event chains. The detailed schema of our database is in the appendix and readers can also refer to the MAGIC-Video paper for details of these indices.

3.2. SQL-based queries

MAGIC-Video uses Personalized Page Rank (PPR) to retrieve nodes in the graph by a query generated by the LLM, with which the agent has no control over the relations

between nodes. We still keep a simple semantic search tool for the model to search all text and images using embeddings. For example, if the agent wants to find who ate the apple yesterday at noon, it likely issues a query like "person ate apple yesterday noon" and hopes the PPR returns relevant nodes. Instead, we formulate this using a relational database. Given the same index structure as MAGIC-Video, the agent can write a SQL query like in Listing 1. Even though searches like this could be replaced with dedicated tools like `search_semantic(object="apple", time_s=time_s, time_end=time_end)`, using SQL as an interface provides a unified interface with more freedom and does not require specific design for each case (for example, if we need to search "the person who ate the apple that is not Luke", we would need to re-design the tool signature). It can also be seen as a form of Programmatic Tool Calling. Such an interface can also handle complex queries using advanced SQL features like JOIN, GROUP BY, and sub-queries when needed.

Our work differs from SAG [11] in that the LLM controls the SQL generation and it can use all SQL features, unlike SAG which only used JOIN as a way to expand the seed nodes.

Listing 1. An example SQL query for finding who ate the apple within a given time window.

```
SELECT *
FROM semantic_triples
WHERE predicate IN ('eat', 'eats', 'ate',
, 'eating')
AND object IN ('apple', 'the_apple', 'an_apple')
AND snapshot_s >= $start_s
AND snapshot_s < $end_s;
```

3.3. Trainable Skills

4. Evaluations

4.1. Datasets

We used the two most popular ULVU datasets in the paper: EgoLifeQA [12] and MM-Lifelong [2]. EgoLifeQA uses multiple-choice questions (MCQ) while MM-Lifelong uses open QA. EgoLifeQA contains one week of recording in the same setting while MM-Lifelong has 3 different settings (day, week, and month), each with its own video. The MM-Lifelong week subset has the same video as EgoLifeQA. We argue that MCQ does not reflect realistic usage since most questions asked by users in the real world are open. Additionally, agents can brute-force verify each option to artificially boost accuracy. We therefore reformulated EgoLifeQA into open QA questions (Open-EgoLifeQA) similar to MM-Lifelong and evaluated both using LLM-as-judge for semantic comparison between the ground-truth and predicted answers. To compare with prior work, we used the

Jack view from EgoLifeQA. We evaluated our untrained VideoSQL on both Open-EgoLifeQA and the MM-Lifelong validation set.

To train the VideoSQL Skill, we need to select an independent training set. Ego-R1 Data [10] is a training set built on top of EgoLifeQA and the authors of Ego-R1 trained their method on this dataset and then tested on EgoLifeQA. MM-Lifelong also has a training split as well. However, these training sets overlap with their corresponding validation/testing sets at the video level. While no questions are leaked, we argue that sharing videos could still inflate performance, or at least fail to test the model's ability for cross-domain generalization. To address this, we trained our Skill on EgoLifeQA and tested it on the more challenging MM-Lifelong (for context, MAGIC-Video reached 67.6 on EgoLifeQA but only 24.5 on MM-Lifelong).

4.2. Baselines

We re-ran the MAGIC-Video [4] as our training-free baseline as we share the same indices. Both MAGIC-Video and our method used Qwen-3.8-27B [8] as the policy model and Qwen3-VL-Embedding-8B [5] as embedding model. We also compared with a couple other baselines directly with their reported numbers. We also selected Trace2Skill [6] and SkillOpt [13] for our skill-trained variant, both built on top of the same VideoSQL harness.

References

- [1] [2510.08191] Training-Free Group Relative Policy Optimization. <https://arxiv.org/abs/2510.08191>. 1
- [2] Guo Chen, Lidong Lu, Yicheng Liu, Liangrui Dong, Lidong Zou, Jixin Lv, Zhenquan Li, Xinyi Mao, Baoqi Pei, Shihao Wang, Zhiqi Li, Karan Sapra, Fuxiao Liu, Yindong Zheng, Yifei Huang, Limin Wang, Zhiding Yu, Andrew Tao, Guilin Liu, and Tong Lu. Towards Multimodal Lifelong Understanding: A Dataset and Agentic Baseline. <https://arxiv.org/abs/2603.05484v1>, 2026. 2
- [3] Kunfeng Chen, Qihuang Zhong, Juhua Liu, and Bo Du. SkillCAT: Contrastive, Assessment-Augmented and Topology-AwareSkill Self-Evolution for LLM Agents. <https://arxiv.org/abs/2606.13317v2>, 2026. 1
- [4] Jiazheng Li, Chi-Hao Wu, Yunze Liu, Kaize Ding, Jundong Li, and Chuxu Zhang. Bridging Modalities, Spanning Time: Structured Memory for Ultra-Long Agentic Video Reasoning, 2026. 1, 2
- [5] Mingxin Li, Yanzhao Zhang, Dingkun Long, Keqin Chen, Sibao Song, Shuai Bai, Zhibo Yang, Pengjun Xie, An Yang, Dayiheng Liu, Jingren Zhou, and Junyang Lin. Qwen3-VL-Embedding and Qwen3-VL-Reranker: A Unified Framework for State-of-the-Art Multimodal Retrieval and Ranking, 2026. 2
- [6] Jingwei Ni, Yihao Liu, Xinpeng Liu, Yutao Sun, Mengyu Zhou, Pengyu Cheng, Dexin Wang, Erchao Zhao, Xiaoxi Jiang, and Guanrun Jiang. Trace2Skill: Distill Trajectory-Local Lessons into Transferable Agent Skills, 2026. 1, 2

- 174 [7] Abdelghny Orogat and Essam Mansour. Is Agent Memory a
175 Database? Rethinking Data Foundations for Long-Term AI
176 Agent Memory, 2026. [1](#)
- 177 [8] Qwen Team. Qwen3.8-Max: A new bar for coding and
178 cowork, 2026. [2](#)
- 179 [9] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao
180 Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li,
181 Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Lim-
182 its of Mathematical Reasoning in Open Language Models,
183 2024. [1](#)
- 184 [10] Shulin Tian, Ruiqi Wang, Hongming Guo, Penghao Wu,
185 Yuhao Dong, Xiuying Wang, Jingkan Yang, Hao Zhang,
186 Hongyuan Zhu, and Ziwei Liu. Ego-R1: Chain-of-Tool-
187 Thought for Ultra-Long Egocentric Video Reasoning, 2025.
188 [2](#)
- 189 [11] Yuchao Wu, Junqin Li, Xingcheng Liang, Yongjie Chen,
190 Yinghao Liang, Linyuan Mo, and Guanxian Li. SAG: SQL-
191 Retrieval Augmented Generation with Query-Time Dynamic
192 Hyperedges, 2026. [1](#), [2](#)
- 193 [12] Jingkan Yang, Shuai Liu, Hongming Guo, Yuhao Dong, Xi-
194 amengwei Zhang, Sicheng Zhang, Pengyun Wang, Zitang
195 Zhou, Binzhu Xie, Ziyue Wang, Bei Ouyang, Zhengyu Lin,
196 Marco Cominelli, Zhongang Cai, Yuanhan Zhang, Peiyuan
197 Zhang, Fangzhou Hong, Joerg Widmer, Francesco Gringoli,
198 Lei Yang, Bo Li, and Ziwei Liu. EgoLife: Towards Egocen-
199 tric Life Assistant, 2026. [2](#)
- 200 [13] Yifan Yang, Ziyang Gong, Weiquan Huang, Qihao Yang, Zi-
201 wei Zhou, Zisu Huang, Yan Li, Xuemei Gao, Qi Dai, Bei
202 Liu, Kai Qiu, Yuqing Yang, Dongdong Chen, Xue Yang, and
203 Chong Luo. SkillOpt: Executive Strategy for Self-Evolving
204 Agent Skills, 2026. [1](#), [2](#)